

## 1. Simple Linear Regression Python script using weight as the independent variable (X) and height as the dependent variable (Y) with 30 sample rows with Standardization included

```
In [19]: # Step 1: Import required libraries
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt

from sklearn.linear_model import LinearRegression
from sklearn.model_selection import train_test_split, cross_val_score
from sklearn.metrics import mean_squared_error, r2_score
from sklearn.preprocessing import StandardScaler

# Step 2: Load dataset
df = pd.read_csv('Regression/1.Simple_Linear_Regression/height-weight.csv')
print("Sample data:\n", df.head(), "\n")

# Step 3: Visualize raw data
plt.figure(figsize=(8, 5))
plt.scatter(df['Weight'], df['Height'], color='blue')
plt.title("Weight vs Height")
plt.xlabel("Weight (kg)")
plt.ylabel("Height (cm)")
plt.grid(True)
plt.show()

# Step 4: Define features (X) and target (y)
X = df[['Weight']] # Features must be 2D
y = df['Height']

# Step 5: Split data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# Step 6: Standardize the features
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)
```

```

print("Scaler Mean:", scaler.mean_)
print("Scaler Scale (Std Dev):", scaler.scale_)

# Step 7: Train Linear Regression model on scaled data
model = LinearRegression()
model.fit(X_train_scaled, y_train)

# Step 8: Print learned parameters
print("\nIntercept:", model.intercept_)
print("Coefficient (Slope):", model.coef_[0])

# Step 9: Make predictions on test set
y_pred = model.predict(X_test_scaled)

# Step 10: Evaluate model performance
mse = mean_squared_error(y_test, y_pred)
r2 = r2_score(y_test, y_pred)

print(f"\nMean Squared Error (MSE): {mse:.2f}")
print(f"Root Mean Squared Error (RMSE): {np.sqrt(mse):.2f}")
print(f"R2 Score: {r2:.2f}")
print(f"Model Accuracy (as R2 %): {r2 * 100:.2f}%")

# Step 11: Compare actual vs predicted values
comparison = pd.DataFrame({
    'Actual Height': y_test.values,
    'Predicted Height': y_pred.round(2)
})
print("\nActual vs Predicted:\n", comparison)

# Step 12: Plot regression line using full standardized dataset
X_scaled_full = scaler.transform(X)
y_pred_full = model.predict(X_scaled_full)

plt.figure(figsize=(8, 5))
plt.scatter(X, y, color='blue', label='Original Data')
plt.plot(X, y_pred_full, color='red', label='Regression Line')
plt.title("Linear Regression (Standardized): Weight vs Height")
plt.xlabel("Weight (kg)")
plt.ylabel("Height (cm)")
plt.legend()

```

```

plt.grid(True)
plt.show()

# Step 13: Predict height for new weight values
new_weights = pd.DataFrame({'Weight': [55, 75, 90]})
new_weights_scaled = scaler.transform(new_weights)
new_heights_pred = model.predict(new_weights_scaled)

new_data_results = pd.DataFrame({
    'Weight (kg)': new_weights['Weight'],
    'Predicted Height (cm)': new_heights_pred.round(2)
})
print("\nPredictions for New Records:\n", new_data_results)

# Step 14: Plot new predictions on regression line
print("\nPlotting new Predictions")
plt.figure(figsize=(8, 5))
plt.scatter(X, y, color='blue', label='Training Data')
plt.plot(X, y_pred_full, color='red', label='Regression Line')
plt.scatter(new_weights, new_heights_pred, color='green', marker='x', s=100, label='New Predictions')
plt.title("New Predictions on Regression Line")
plt.xlabel("Weight (kg)")
plt.ylabel("Height (cm)")
plt.legend()
plt.grid(True)
plt.show()

# Step 15: Residual plot (to check model fit)
residuals = y_test - y_pred
plt.figure(figsize=(8, 5))
plt.scatter(y_pred, residuals, color='purple')
plt.axhline(y=0, color='black', linestyle='--')
plt.title("Residual Plot")
plt.xlabel("Predicted Height")
plt.ylabel("Residuals (Actual - Predicted)")
plt.grid(True)
plt.show()

# Step 16: Cross-validation (5-fold)
cv_scores = cross_val_score(model, X_train_scaled, y_train, cv=5, scoring='neg_mean_squared_error')
cv_scores = -cv_scores # Convert from negative MSE to positive

```

```

cv_mean = np.mean(cv_scores)
cv_std = np.std(cv_scores)

print(f"\nCross-validation MSE (Mean): {cv_mean:.2f}")
print(f"Cross-validation MSE (Std Dev): {cv_std:.2f}")
print(f"Cross-validation RMSE (Mean): {np.sqrt(cv_mean):.2f}")


# Step 17: Compute Adjusted R2
n = len(X_train) # Number of samples
p = X_train.shape[1] # Number of features
adj_r2 = 1 - ((1 - r2) * (n - 1)) / (n - p - 1)
print(f"\nAdjusted R2: {adj_r2:.2f}")

```

Sample data:

|   | Weight | Height |
|---|--------|--------|
| 0 | 45     | 120    |
| 1 | 58     | 135    |
| 2 | 48     | 123    |
| 3 | 60     | 145    |
| 4 | 70     | 160    |

Weight vs Height




Scaler Mean: [74.27777778]  
Scaler Scale (Std Dev): [17.68805484]

Intercept: 157.5  
Coefficient (Slope): 17.034408719095538

Mean Squared Error (MSE): 109.78  
Root Mean Squared Error (RMSE): 10.48  
R<sup>2</sup> Score: 0.78  
Model Accuracy (as R<sup>2</sup> %): 77.70%

Actual vs Predicted:

|   | Actual Height | Predicted Height |
|---|---------------|------------------|
| 0 | 177           | 161.08           |
| 1 | 170           | 161.08           |
| 2 | 120           | 129.30           |
| 3 | 182           | 177.46           |
| 4 | 159           | 148.57           |

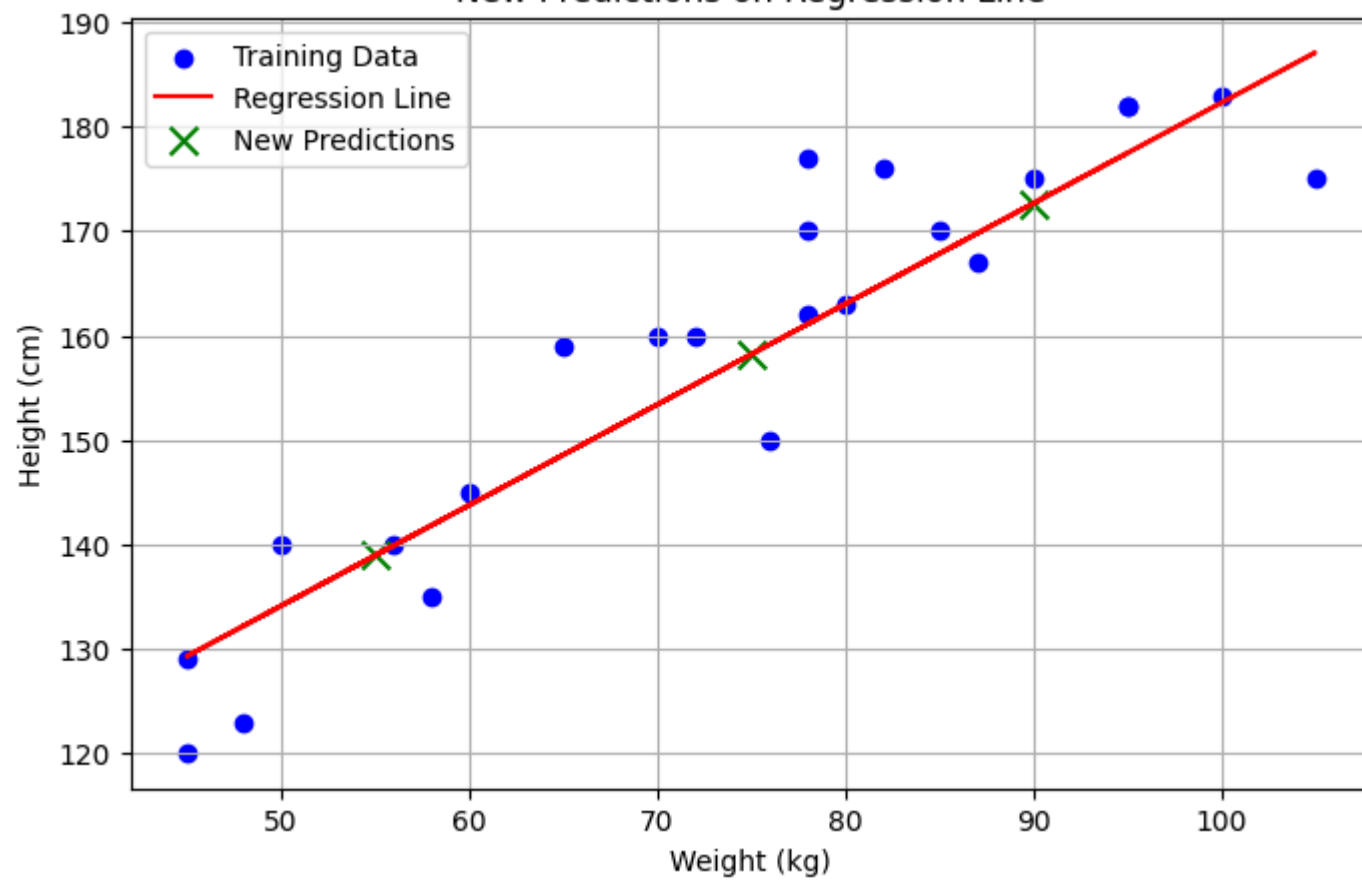


Predictions for New Records:


|   | Weight (kg) | Predicted Height (cm) |
|---|-------------|-----------------------|
| 0 | 55          | 138.93                |
| 1 | 75          | 158.20                |
| 2 | 90          | 172.64                |

Plotting new Predictions

New Predictions on Regression Line







Cross-validation MSE (Mean): 50.23  
Cross-validation MSE (Std Dev): 25.78  
Cross-validation RMSE (Mean): 7.09

Adjusted  $R^2$ : 0.76

2. Multiple Linear Regression example in Python using 2 features — Weight(kg) and Age — to predict Height(cm), with standardization, evaluation, and inline comments for every major step:

```
In [20]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
```

```

from sklearn.linear_model import LinearRegression
from sklearn.model_selection import train_test_split, cross_val_score
from sklearn.preprocessing import StandardScaler
from sklearn.metrics import mean_squared_error, r2_score

# 1. Generate synthetic data
np.random.seed(42)
weights = np.random.normal(loc=65, scale=15, size=30).round(2) # Weight in kg
ages = np.random.randint(18, 60, size=30) # Age in years
# Height depends on weight and age (with some noise)
heights = (weights * 0.7 + ages * 0.2 + np.random.normal(0, 4, 30)).round(2) + 90

# 2. Create a DataFrame
df = pd.DataFrame({
    'Weight(kg)': weights,
    'Age': ages,
    'Height(cm)': heights
})

print("Sample data:\n", df.head(), "\n")

# 3. Define features (X) and target (y)
X = df[['Weight(kg)', 'Age']] # Two features
y = df['Height(cm)']

# 4. Train-test split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# 5. Standardize the features
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)

print("Feature means after standardization (should be ~0):", X_train_scaled.mean(axis=0))
print("Feature std devs (should be ~1):", X_train_scaled.std(axis=0), "\n")

# 6. Train the Linear Regression model
model = LinearRegression()
model.fit(X_train_scaled, y_train)

```

```

# 7. Output coefficients
print("Intercept (bias term):", model.intercept_)
print("Coefficients (slopes):")
for feature, coef in zip(X.columns, model.coef_):
    print(f" {feature}: {coef:.4f}")

# 8. Predict using test data
y_pred = model.predict(X_test_scaled)

# 9. Evaluate the model
mse = mean_squared_error(y_test, y_pred)
r2 = r2_score(y_test, y_pred)

print(f"\nMean Squared Error: {mse:.2f}")
print(f"R² Score: {r2:.2f}")

# 10. Cross-validation
cv_scores = cross_val_score(model, X_train_scaled, y_train, cv=5, scoring='neg_mean_squared_error')
cv_scores = -cv_scores # Convert negative MSE to positive
cv_mean = np.mean(cv_scores)
cv_std = np.std(cv_scores)
print(f"\nCross-validation MSE (Mean): {cv_mean:.2f}")
print(f"Cross-validation RMSE (Mean): {np.sqrt(cv_mean):.2f}")

# 11. Compare actual vs predicted
results = pd.DataFrame({
    'Weight': X_test['Weight(kg)'].values,
    'Age': X_test['Age'].values,
    'Actual Height': y_test.values,
    'Predicted Height': y_pred.round(2)
})

print("\nActual vs Predicted:\n", results)

# 12. Residual plot (for visualizing model fit)
residuals = y_test - y_pred

plt.figure(figsize=(8, 5))
plt.scatter(y_pred, residuals, color='purple')
plt.axhline(y=0, color='red', linestyle='--')
plt.title("Residuals vs Predicted Values")

```

```

plt.xlabel("Predicted Height (cm)")
plt.ylabel("Residuals (Actual - Predicted)")
plt.grid(True)
plt.show()

# 13. Test for new records (Example: new weights and ages)
new_data = pd.DataFrame({
    'Weight(kg)': [70, 80, 90], # New data points
    'Age': [25, 30, 35]
})

# Standardize the new data
new_data_scaled = scaler.transform(new_data)

# Predict heights for new records
new_predictions = model.predict(new_data_scaled)

# Display new predictions
new_results = pd.DataFrame({
    'Weight(kg)': new_data['Weight(kg)'],
    'Age': new_data['Age'],
    'Predicted Height(cm)': new_predictions.round(2)
})

print("\nPredictions for New Records:\n", new_results)

# 14. Plotting new predictions along with original data and regression line
plt.figure(figsize=(8, 5))
plt.scatter(X['Weight(kg)'], y, color='blue', label='Training Data')
plt.scatter(new_data['Weight(kg)'], new_predictions, color='green', marker='x', s=100, label='New Predictions')
plt.title("New Predictions on Regression Line")
plt.xlabel("Weight (kg)")
plt.ylabel("Height (cm)")
plt.legend()
plt.grid(True)
plt.show()

```

Sample data:

|   | Weight(kg) | Age | Height(cm) |
|---|------------|-----|------------|
| 0 | 72.45      | 37  | 157.48     |
| 1 | 62.93      | 45  | 147.33     |
| 2 | 74.72      | 24  | 147.49     |
| 3 | 87.85      | 25  | 158.17     |
| 4 | 61.49      | 52  | 139.63     |

Feature means after standardization (should be ~0): [-2.59052039e-16 -2.22044605e-16]

Feature std devs (should be ~1): [1. 1.]

Intercept (bias term): 141.54625000000001

Coefficients (slopes):

Weight(kg): 8.8222

Age: 1.1855

Mean Squared Error: 12.10


R<sup>2</sup> Score: 0.75

Cross-validation MSE (Mean): 21.47

Cross-validation RMSE (Mean): 4.63


Actual vs Predicted:

|   | Weight | Age | Actual Height | Predicted Height |
|---|--------|-----|---------------|------------------|
| 0 | 70.64  | 38  | 143.67        | 146.93           |
| 1 | 56.57  | 35  | 136.15        | 137.70           |
| 2 | 43.63  | 25  | 125.22        | 128.48           |
| 3 | 69.71  | 51  | 144.50        | 147.66           |
| 4 | 57.96  | 57  | 144.52        | 140.82           |
| 5 | 73.14  | 21  | 141.77        | 146.79           |



Predictions for New Records:

|   | Weight(kg) | Age | Predicted Height(cm) |
|---|------------|-----|----------------------|
| 0 | 70         | 25  | 145.21               |
| 1 | 80         | 30  | 152.06               |
| 2 | 90         | 35  | 158.91               |



3. Multiple Linear Regression example to include a categorical variable — for instance, Gender — and walk through how to encode it and use it in the regression model.

```
In [21]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt

from sklearn.linear_model import LinearRegression
from sklearn.model_selection import train_test_split, cross_val_score
from sklearn.preprocessing import StandardScaler
from sklearn.metrics import mean_squared_error, r2_score
```

```

# 1. Generate synthetic data
np.random.seed(42)
weights = np.random.normal(loc=65, scale=15, size=30).round(2)
ages = np.random.randint(18, 60, size=30)
genders = np.random.choice(['Male', 'Female'], size=30)

# Simulate height: depend on weight, age, and gender (with noise)
# Let's assume 'Male' adds about 5 cm more on average
gender_bias = np.where(genders == 'Male', 5, 0)
heights = (weights * 0.7 + ages * 0.2 + gender_bias + np.random.normal(0, 4, 30)).round(2) + 90

# 2. Create DataFrame
df = pd.DataFrame({
    'Weight(kg)': weights,
    'Age': ages,
    'Gender': genders,
    'Height(cm)': heights
})

print("Sample data with Gender:\n", df.head(), "\n")

# 3. One-hot encode the categorical variable 'Gender'
df_encoded = pd.get_dummies(df, columns=['Gender'], drop_first=True)
# 'Gender_Male' will be 1 if Male, 0 if Female

# 4. Separate features and target
X = df_encoded.drop(columns=['Height(cm)'])
y = df_encoded['Height(cm)']

# 5. Train-test split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# 6. Standardize only numeric columns (not dummy variables)
numeric_features = ['Weight(kg)', 'Age']
scaler = StandardScaler()

# Fit on training numeric columns
X_train_scaled = X_train.copy()
X_test_scaled = X_test.copy()

```



```

X_train_scaled[numeric_features] = scaler.fit_transform(X_train[numeric_features])
X_test_scaled[numeric_features] = scaler.transform(X_test[numeric_features])

# 7. Train the regression model
model = LinearRegression()
model.fit(X_train_scaled, y_train)

# 8. Output coefficients
print("Intercept (bias):", model.intercept_)
print("Coefficients:")
for feature, coef in zip(X.columns, model.coef_):
    print(f" {feature}: {coef:.4f}")

# 9. Predict and evaluate
y_pred = model.predict(X_test_scaled)
mse = mean_squared_error(y_test, y_pred)
r2 = r2_score(y_test, y_pred)

print(f"\nMean Squared Error: {mse:.2f}")
print(f"R2 Score: {r2:.2f}")

# 10. Cross-validation
cv_scores = cross_val_score(model, X_train_scaled, y_train, cv=5, scoring='neg_mean_squared_error')
cv_scores = -cv_scores # Convert negative MSE to positive
cv_mean = np.mean(cv_scores)
cv_std = np.std(cv_scores)
print(f"\nCross-validation MSE (Mean): {cv_mean:.2f}")
print(f"Cross-validation RMSE (Mean): {np.sqrt(cv_mean):.2f}")

# 11. Compare actual vs predicted
results = pd.DataFrame({
    'Weight': X_test['Weight(kg)'].values,
    'Age': X_test['Age'].values,
    'Gender_Male': X_test['Gender_Male'].values,
    'Actual Height': y_test.values,
    'Predicted Height': y_pred.round(2)
})

print("\nActual vs Predicted with Gender:\n", results)

# 12. Residual plot

```

```

residuals = y_test - y_pred
plt.figure(figsize=(8, 5))
plt.scatter(y_pred, residuals, color='green')
plt.axhline(y=0, color='red', linestyle='--')
plt.title("Residuals vs Predicted (with Gender)")
plt.xlabel("Predicted Height (cm)")
plt.ylabel("Residuals")
plt.grid(True)
plt.show()

# 13. Test for new records (Example: new weights, ages, and genders)
new_data = pd.DataFrame({
    'Weight(kg)': [70, 80, 90], # New data points
    'Age': [25, 30, 35],
    'Gender': ['Male', 'Female', 'Male']
})

# One-hot encode new gender data
new_data_encoded = pd.get_dummies(new_data, columns=['Gender'], drop_first=True)

# Standardize the new data (same scaler)
new_data_scaled = new_data_encoded.copy()
new_data_scaled[numeric_features] = scaler.transform(new_data_encoded[numeric_features])

# Predict heights for new records
new_predictions = model.predict(new_data_scaled)

# Display new predictions
new_results = pd.DataFrame({
    'Weight(kg)': new_data['Weight(kg)'],
    'Age': new_data['Age'],
    'Gender_Male': new_data_encoded['Gender_Male'],
    'Predicted Height(cm)': new_predictions.round(2)
})

print("\nPredictions for New Records:\n", new_results)

# 14. Plotting new predictions along with original data and regression line
plt.figure(figsize=(8, 5))
plt.scatter(X['Weight(kg)'], y, color='blue', label='Training Data')
plt.scatter(new_data['Weight(kg)'], new_predictions, color='green', marker='x', s=100, label='New Predictions')

```

```
plt.title("New Predictions on Regression Line")
plt.xlabel("Weight (kg)")
plt.ylabel("Height (cm)")
plt.legend()
plt.grid(True)
plt.show()
```

Sample data with Gender:

|   | Weight(kg) | Age | Gender | Height(cm) |
|---|------------|-----|--------|------------|
| 0 | 72.45      | 37  | Male   | 155.91     |
| 1 | 62.93      | 45  | Male   | 148.75     |
| 2 | 74.72      | 24  | Female | 149.75     |
| 3 | 87.85      | 25  | Female | 157.46     |
| 4 | 61.49      | 52  | Male   | 147.77     |

Intercept (bias): 141.506117273088

Coefficients:

Weight(kg): 9.6636  
 Age: 0.6915  
 Gender\_Male: 3.7302

Mean Squared Error: 18.03

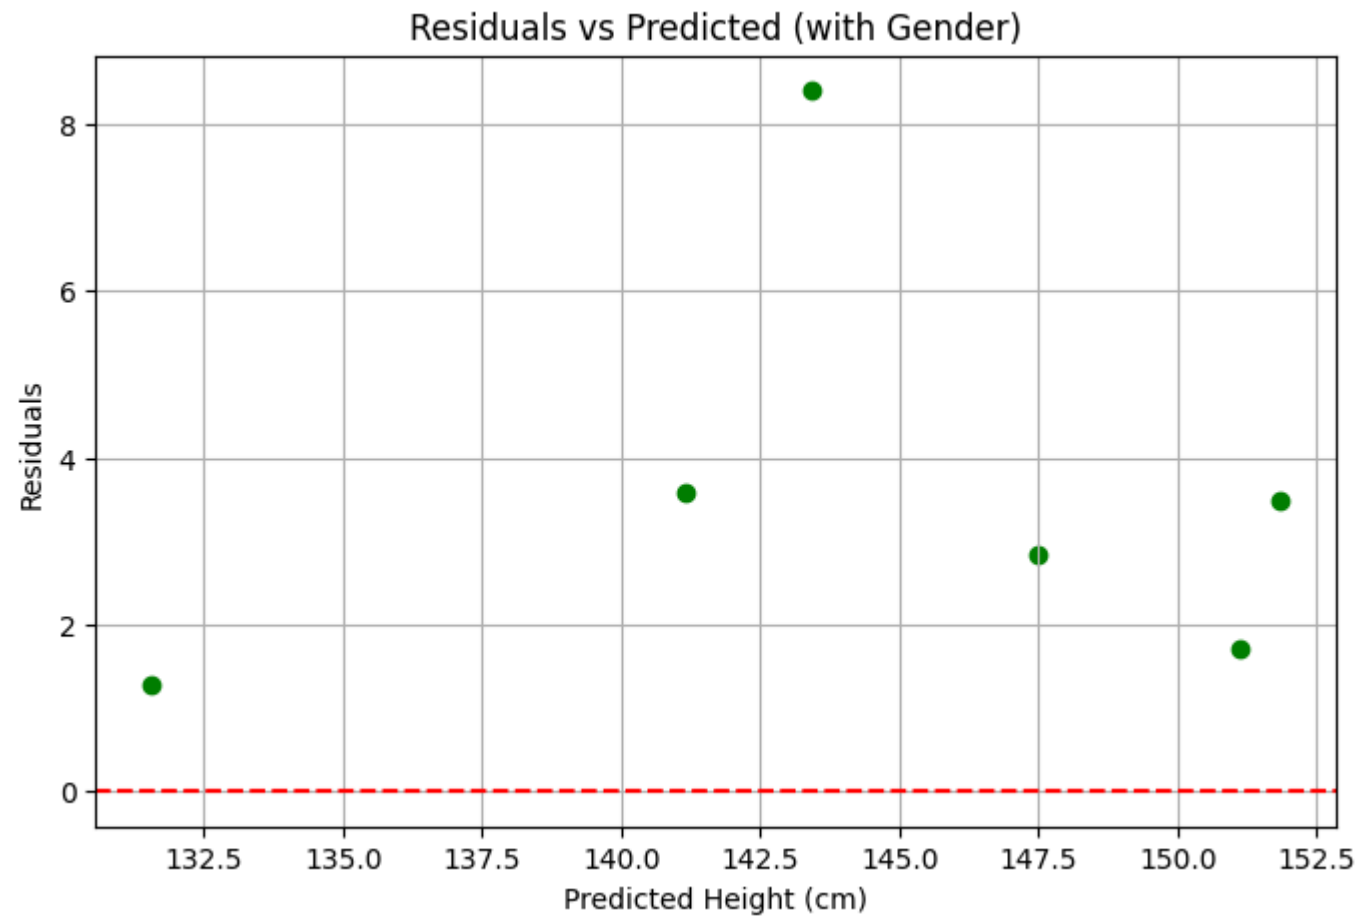
R<sup>2</sup> Score: 0.68

Cross-validation MSE (Mean): 18.19

Cross-validation RMSE (Mean): 4.27


Actual vs Predicted with Gender:

|   | Weight | Age | Gender_Male | Actual Height | Predicted Height |
|---|--------|-----|-------------|---------------|------------------|
| 0 | 70.64  | 38  | True        | 152.81        | 151.11           |
| 1 | 56.57  | 35  | True        | 144.74        | 141.15           |
| 2 | 43.63  | 25  | True        | 132.86        | 131.57           |
| 3 | 69.71  | 51  | False       | 150.33        | 147.50           |
| 4 | 57.96  | 57  | True        | 151.82        | 143.42           |
| 5 | 73.14  | 21  | True        | 155.33        | 151.84           |



Predictions for New Records:

|   | Weight(kg) | Age | Gender_Male | Predicted Height(cm) |
|---|------------|-----|-------------|----------------------|
| 0 | 70         | 25  | True        | 149.89               |
| 1 | 80         | 30  | False       | 153.41               |
| 2 | 90         | 35  | True        | 164.38               |



## 4. Polynomial Regression:

- ✓ Generates synthetic data with gender, weight, and age
- ✓ Applies polynomial feature transformation
- ✓ Splits and standardizes data
- ✓ Trains a model and evaluates it
- ✓ Predicts height for 10 new samples at the end

```
In [22]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt

from sklearn.linear_model import LinearRegression
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler, PolynomialFeatures
from sklearn.metrics import mean_squared_error, r2_score

# 1. Generate synthetic data
np.random.seed(42)
weights = np.random.normal(loc=65, scale=15, size=30).round(2)
ages = np.random.randint(18, 60, size=30)
genders = np.random.choice(['Male', 'Female'], size=30)

# Add gender influence: +5 cm for male
gender_bias = np.where(genders == 'Male', 5, 0)
heights = (weights * 0.7 + ages * 0.2 + gender_bias + np.random.normal(0, 4, 30)).round(2) + 90

# 2. Create DataFrame
df = pd.DataFrame({
    'Weight(kg)': weights,
    'Age': ages,
    'Gender': genders,
    'Height(cm)': heights
})

print("Sample data:\n", df.head(), "\n")

# 3. One-hot encode Gender
df_encoded = pd.get_dummies(df, columns=['Gender'], drop_first=True)

# 4. Feature-target split
X = df_encoded.drop(columns=['Height(cm)'])
y = df_encoded['Height(cm)']

# 5. Polynomial Feature Transformation (degree 2)
poly = PolynomialFeatures(degree=2, include_bias=False)
X_poly = poly.fit_transform(X)
```

```

# 6. Train-test split
X_train, X_test, y_train, y_test = train_test_split(X_poly, y, test_size=0.2, random_state=42)

# 7. Standardize
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)

# 8. Train the model
model = LinearRegression()
model.fit(X_train_scaled, y_train)

# 9. Coefficients
print("\nIntercept:", model.intercept_)
print("Coefficients:")
for name, coef in zip(poly.get_feature_names_out(X.columns), model.coef_):
    print(f" {name}: {coef:.4f}")

# 10. Predict and evaluate
y_pred = model.predict(X_test_scaled)
mse = mean_squared_error(y_test, y_pred)
r2 = r2_score(y_test, y_pred)
print(f"\nMean Squared Error: {mse:.2f}")
print(f"R2 Score: {r2:.2f}")

# 11. Actual vs Predicted
X_test_df = pd.DataFrame(X_test, columns=poly.get_feature_names_out(X.columns))
results = pd.DataFrame({
    'Actual Height': y_test.values,
    'Predicted Height': y_pred.round(2)
})
print("\nActual vs Predicted:\n", results)

# 12. Residual plot
residuals = y_test - y_pred
plt.figure(figsize=(8, 5))
plt.scatter(y_pred, residuals, color='orange')
plt.axhline(y=0, color='red', linestyle='--')
plt.title("Residuals vs Predicted")
plt.xlabel("Predicted Height (cm)")
plt.ylabel("Residuals")

```

```
plt.grid(True)
plt.show()

# 13. Predicting multiple new records
new_samples = pd.DataFrame({
    'Weight(kg)': [72, 60, 68, 55, 80, 77, 62, 59, 70, 65],
    'Age': [25, 32, 40, 22, 28, 35, 30, 27, 45, 38],
    'Gender_Male': [1, 0, 1, 0, 1, 1, 0, 0, 1, 0]
})

# Transform and predict
new_samples_poly = poly.transform(new_samples)
new_samples_scaled = scaler.transform(new_samples_poly)
new_preds = model.predict(new_samples_scaled)

# Combine input and prediction
new_results = new_samples.copy()
new_results['Predicted Height (cm)'] = new_preds.round(2)

print("\nPredicted Heights for New Samples:\n", new_results)
```



Sample data:

|   | Weight(kg) | Age | Gender | Height(cm) |
|---|------------|-----|--------|------------|
| 0 | 72.45      | 37  | Male   | 155.91     |
| 1 | 62.93      | 45  | Male   | 148.75     |
| 2 | 74.72      | 24  | Female | 149.75     |
| 3 | 87.85      | 25  | Female | 157.46     |
| 4 | 61.49      | 52  | Male   | 147.77     |

Intercept: 143.8375

Coefficients:

Weight(kg): -2.4436

Age: 7.1496

Gender\_Male: -5.8697

Weight(kg)^2: 4.4067

Weight(kg) Age: 7.6409

Weight(kg) Gender\_Male: 8.8809

Age^2: -15.3730

Age Gender\_Male: 3.1451

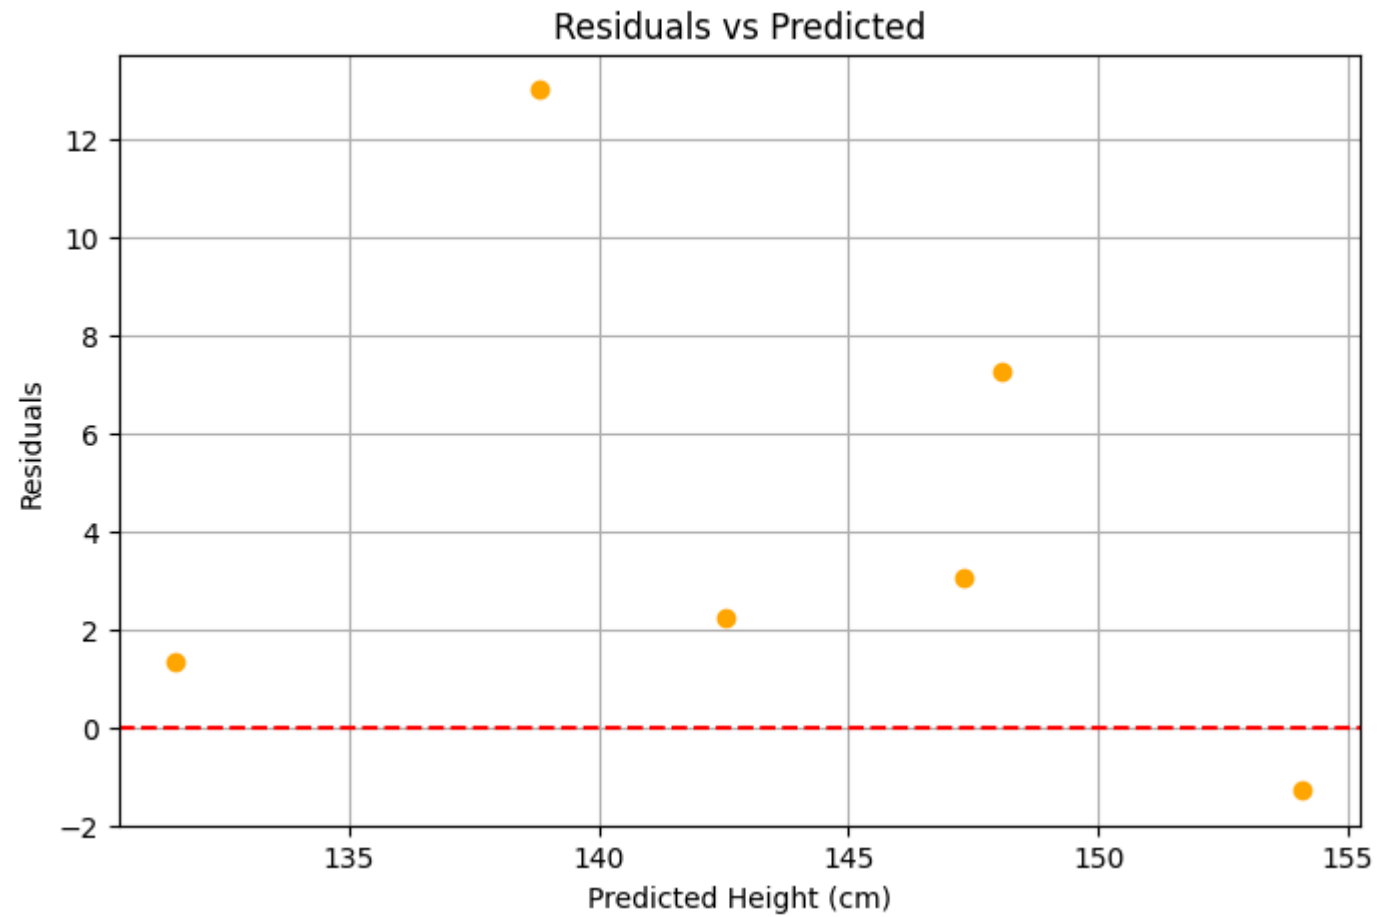
Gender\_Male^2: -5.8697

Mean Squared Error: 40.00

R<sup>2</sup> Score: 0.29

Actual vs Predicted:

|   | Actual Height | Predicted Height |
|---|---------------|------------------|
| 0 | 152.81        | 154.10           |
| 1 | 144.74        | 142.52           |
| 2 | 132.86        | 131.51           |
| 3 | 150.33        | 147.29           |
| 4 | 151.82        | 138.80           |
| 5 | 155.33        | 148.06           |



Predicted Heights for New Samples:

|   | Weight(kg) | Age | Gender_Male | Predicted Height (cm) |
|---|------------|-----|-------------|-----------------------|
| 0 | 72         | 25  | 1           | 150.02                |
| 1 | 60         | 32  | 0           | 145.13                |
| 2 | 68         | 40  | 1           | 152.13                |
| 3 | 55         | 22  | 0           | 140.78                |
| 4 | 80         | 28  | 1           | 157.90                |
| 5 | 77         | 35  | 1           | 158.69                |
| 6 | 62         | 30  | 0           | 145.66                |
| 7 | 59         | 27  | 0           | 143.79                |
| 8 | 70         | 45  | 1           | 153.90                |
| 9 | 65         | 38  | 0           | 147.66                |